



# MACHINE LEARNING

## LINEAR REGRESSION

# AGENDA

## **01** Supervised Learning

Training data, hypotheses, an example.

## **02** LMS Algorithm

Linear model, cost function, gradient descent.

## **03** Normal Equations

Matrix linear model, deriving the normal equations

## **04** Probabilistic Interpretation

Assumptions for errors, likelihood function

## **05** Base Function

Derivation, functions, maximum likelihood





# **SUPERVISED LEARNING**

DEFINITIONS

NOMENCLATURE

# SUPERVISED LEARNING

## DEFINITIONS AND NOMENCLATURE




We start with a data set that represents first-floor living space and sales prices for 1,460 homes in Ames, Iowa.


1 <sup>st</sup> floor square feet	Sale Price
856	208500
1262	181500
920	223500
961	140000
1145	250000
⋮	⋮

# SUPERVISED LEARNING

## DEFINITIONS AND NOMENCLATURE



We graph the 1,460 elements using Python:



# SUPERVISED LEARNING

## DEFINITIONS AND NOMENCLATURE



The question that naturally arises would be:

**HOW DO WE PREDICT THE PRICES FOR OTHER  
HOMES IN AMES, IOWA?**

# SUPERVISED LEARNING

## DEFINITIONS AND NOMENCLATURE



To answer the question, we must first define the nomenclature that will be used:

Variable / Symbol	Description	Example
$x^{(i)}$	Input variable or characteristics	The house's first floor area
$y^{(i)}$	Output, target, or response variable	The house's sale price
$(x^{(i)}, y^{(i)})$	Training example	(area, price)
$m$	Number of training examples	1460 houses with their first floor area and their price.
$\{(x^{(i)}, y^{(i)}); i = 1, \dots, m\}$	Training data set	NA
$\mathcal{X}$	Input value space	NA
$\mathcal{Y}$	Output value space	

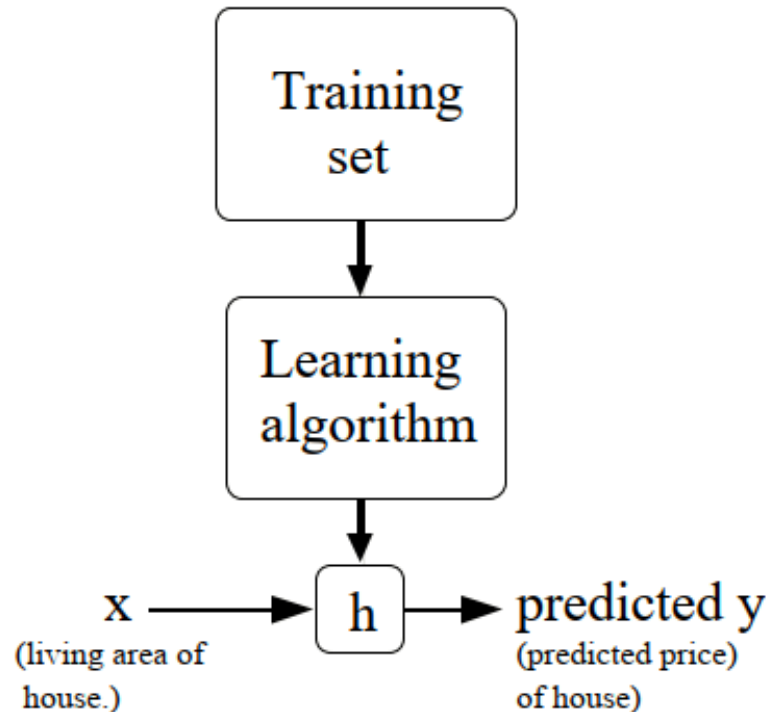
# SUPERVISED LEARNING

## DEFINITIONS AND NOMENCLATURE



The **main goal** of any **supervised algorithm** is to **learn** a **function**  $h: \mathcal{X} \rightarrow \mathcal{Y}$ , such that  $h(x)$  can **predict** the corresponding value of  $y$ .

Where  $h$  is defined as the **hypothesis**.





# SUPERVISED LEARNING

## DEFINITIONS AND NOMENCLATURE



There are **2 types** of **supervised learning problems**:

Type of problem	Description
Regression problem	$y$ takes continuous values.
Classification problem	$y$ takes discrete values



# LMS ALGORITHM

LINEAR MODEL

COST FUNCTION

GRADIENT DESCENT

# L M S A L G O R I T H M

## L I N E A R M O D E L



Now let's look at the same model but with another variable: the living space on the second floor.

1 <sup>st</sup> floor square feet	2 <sup>nd</sup> floor square feet	Sale Price
856	854	208500
1262	0	181500
920	866	223500
961	756	140000
1145	1053	250000
⋮	⋮	⋮


In this case each  $x^{(i)} \in \mathbb{R}^2$ , so it is defined as a vector  $x^{(i)} = [x^{(i)}_1 \ x^{(i)}_2]$ .

# LS ALGORITHM

## LINEAR MODEL




We graph the 1,460 elements again using Python:



# L M S A L G O R I T H M

## L I N E A R M O D E L



It is **hypothesized** that the data is distributed in a **linear fashion**, therefore:

$$h_w(x) = w_0 + w_1x_1 + w_2x_2$$


In this case the **variables**  $w_i$  are defined as the **parameters** or **weights** that **parameterize** the **space** of all **linear functions** that map from  $\mathcal{X}$  to  $\mathcal{Y}$ .

The expression is **simplified** into **vector form**, where  $n$  represents the **number of input variables** (omitting  $x_0$ ),:

$$h(x) = \sum_{i=1}^n w_i x_i = w^T x$$

# L M S    A L G O R I T H M

## C O S T    F U N C T I O N



Then, the question would be:

**WHAT IS THE BEST COMBINATION OF PARAMETERS  $w$   
THAT RESULTS IN THE BEST HYPOTHESIS  $h$ ?**

# L M S A L G O R I T H M

## C O S T F U N C T I O N



To **answer** the **question**, we need to **measure** the **error** produced by the **estimates** produced by  $h$ .

We know that we can **measure** the **error** of a single **estimate**  $i$  as a **difference** **between** the estimate  $h_w(x^{(i)})$  and the response variable  $y^{(i)}$ .


$$e = h_w(x^{(i)}) - y^{(i)}$$

To ensure positive amounts of error, the answer is squared:


$$e^2 = (h_w(x^{(i)}) - y^{(i)})^2$$

# LEAST SQUARES ALGORITHM

## COST FUNCTION



The **error** that we are **calculating graphically** is **interpreted** with an **example**:




The linear model  
represented by  $h$

Response variables  $y^{(i)}$



# L M S A L G O R I T H M

## C O S T F U N C T I O N



So far, we have calculated the **squared error** of a **single training data**. The **mean square error** *MSE* is calculated for **all training data**.


$$MSE = \frac{1}{2m} \sum_{i=1}^m (h_w(x^{(i)}) - y^{(i)})^2$$

$$MSE = J(w)$$

Where  $J(w)$  is the **cost function**.

# L M S    A L G O R I T H M

## G R A D I E N T    D E S C E N T



Therefore, the **best combination** of weights  $w$  is the one that **minimizes** the **cost function**  $J(w)$ , which **measures** our **mean square error**.


To find the **best combination**, let's design a **search algorithm** that **starts** with a **random initial value** of  $w$ , and **updates** the **values** of  $w$  until it **converges** to a **minimum value** of  $J(w)$ . The **constant**  $\alpha$  is defined as the **learning rate**.

$$w_j := w_j - \alpha \frac{\partial}{\partial w_j} J(w)$$


**NOTE:** the above equation **only updates** the **value** of a **single weight**  $w_j$  of the  $n$  weights that **parameterize** the **linear model**. In reality, **all weights**  $w_j$  are **updated simultaneously**.

# L M S     A L G O R I T H M

## G R A D I E N T     D E S C E N T




In the **gradient descent optimization algorithm**, we **update the weights** in the **direction** with the **greatest decrement of  $J(w)$** .



# L M S A L G O R I T H M

## G R A D I E N T D E S C E N T



The **derivative** of the **cost function**  $J(\mathbf{w})$  with respect to a **specific weight**  $w_j$  is calculated.

**Derive** the result:

$$\frac{\partial}{\partial w_j} J(\mathbf{w}) = \frac{1}{m} (h(\mathbf{x}) - \mathbf{y}) x_j$$

# L M S     A L G O R I T H M

## G R A D I E N T     D E S C E N T



Therefore, the **gradient descent weights update** for a **single training data** would look like this:

$$w_j := w_j + \frac{\alpha}{m} (y^{(i)} - h(x^{(i)})) x_j$$

This **equation** is called the **LMS update rule** (“**Least Mean Squares**”) or also the **Widrow-Hoff learning rule**.

This **method** is also called **batch gradient descent**, where **all training data** is **analyzed simultaneously**.

# L M S     A L G O R I T H M

## G R A D I E N T     D E S C E N T



For  $m$  training data, and defining  $x^{(i)}$  and  $w$  as **vectors**, the learning rule would look like this:

$$w := w + \frac{\alpha}{m} \sum_{i=1}^m (y^{(i)} - h(x^{(i)})) x^{(i)}$$

This **equation** is called the **LMS update rule** (“Least Mean Squares”) or also the **Widrow-Hoff learning rule**.


# L M S A L G O R I T H M

## G R A D I E N T D E S C E N T




In general, the **gradient descent method** can suffer from **several local minima**. In this case, for **linear regression models** there is only a **single global minimum**.

Consequently, the **algorithm will always converge** assuming that the **learning rate is not too high**.



# L M S    A L G O R I T H M

## G R A D I E N T    D E S C E N T



When the **method** only **observes one training data** at a time, and **updates the weights with a single data**, it is called **stochastic or incremental gradient descent**.


$$\mathbf{w} := \mathbf{w} + \alpha(\mathbf{y}^{(i)} - \mathbf{h}(\mathbf{x}^{(i)}))\mathbf{x}^{(i)}$$

This **variant** of the algorithm is **used** when it is **very expensive to evaluate** the **update** for **very large data sets** (when  $m$  is **very large**), but it has the minimal **divergence** problem.



# LEAST SQUARES ALGORITHM

## GRADIENT DESCENT




### Example:

Random training data was generated with a certain degree of error  $e$ :

$$y = 3x + 4 + e$$

and both weights were randomly initialized:  $w_0$  and  $w_1$




# L M S A L G O R I T H M

## G R A D I E N T D E S C E N T



### Example:

After 1000 iterations the following model was obtained:



$$w_0 = 4.2174$$
$$w_1 = 2.816$$

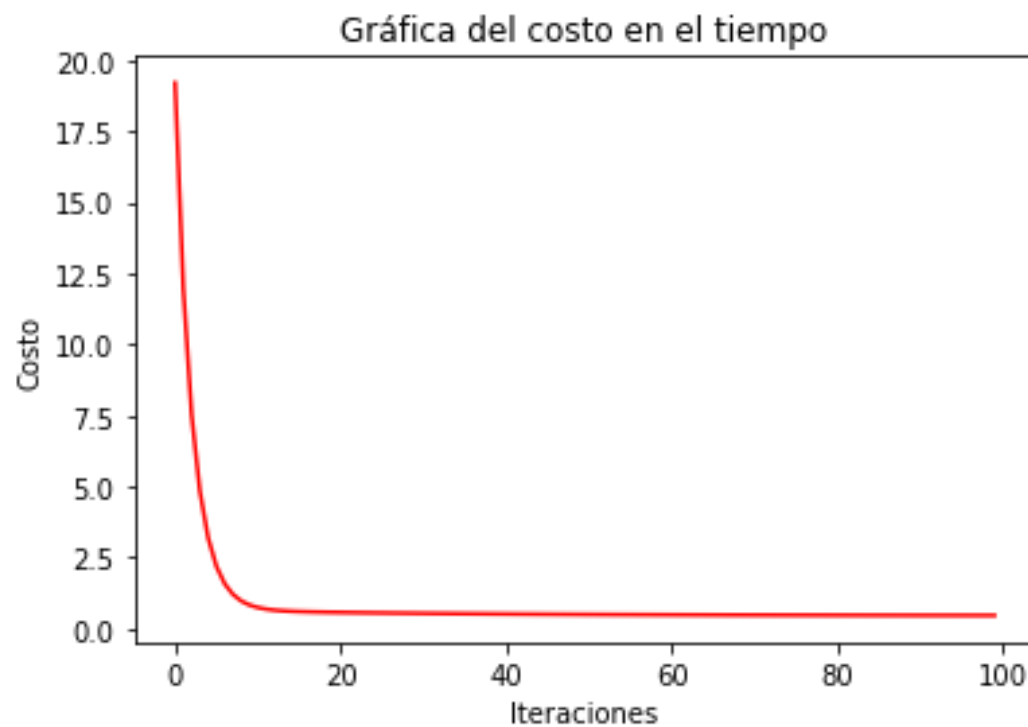
# L M S A L G O R I T H M

## G R A D I E N T D E S C E N T



**Example:**

Error based on 100 iterations.





# **NORMAL EQUATIONS**

LINEAR MATRIX MODEL  
DERIVATION

# NORMAL EQUATIONS

## MATRIX NOTATION



Now we want to **calculate** the **minimum** of the **cost function analytically**. For this, **matrix notation** is introduced, which **allows** the **derivatives** of the **cost function** to be **expressed** in an **elegant way** without getting into so much **verbiage**.

The  **$m$  training data**  $\{(x^{(i)}, y^{(i)}); i = 1, \dots, m\}$  is represented as a matrix  $X \in \mathbb{R}^{m \times n}$ , the set of **outputs** as a **vector**  $\vec{y} \in \mathbb{R}^m$  and the  **$n$  weights**  $w_j$  as a vector  $\vec{w} \in \mathbb{R}^n$ .

$$x^{(i)} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \quad X = \begin{bmatrix} -(x^{(1)})^T & - \\ -(x^{(2)})^T & - \\ \vdots & \\ -(x^{(m)})^T & - \end{bmatrix} \quad \vec{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix} \quad \vec{w} = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_n \end{bmatrix}$$

# NORMAL EQUATIONS

## MATRIX NOTATION




The **absolute error function** is represented in **matrix notation** where  $\mathbf{h}_w(\mathbf{x}^{(i)}) = (\mathbf{x}^{(i)})^T \vec{w}$ :

$$\mathbf{X}\vec{w} - \vec{y} = \begin{bmatrix} -(\mathbf{x}^{(1)})^T \vec{w} - \\ -(\mathbf{x}^{(2)})^T \vec{w} - \\ \vdots \\ -(\mathbf{x}^{(m)})^T \vec{w} - \end{bmatrix} - \begin{bmatrix} \mathbf{y}^{(1)} \\ \mathbf{y}^{(2)} \\ \vdots \\ \mathbf{y}^{(m)} \end{bmatrix}$$

$$\mathbf{X}\vec{w} - \vec{y} = \begin{bmatrix} \mathbf{h}_w(\mathbf{x}^{(1)}) - \mathbf{y}^{(1)} \\ \mathbf{h}_w(\mathbf{x}^{(2)}) - \mathbf{y}^{(2)} \\ \vdots \\ \mathbf{h}_w(\mathbf{x}^{(m)}) - \mathbf{y}^{(m)} \end{bmatrix}$$

# NORMAL EQUATIONS

## MATRIX NOTATION



The **mean square error term** (cost function):


$$J(\mathbf{w}) = \frac{1}{2m} \sum_{i=1}^m (h_{\mathbf{w}}(\mathbf{x}^{(i)}) - \mathbf{y}^{(i)})^2$$

can be **represented** in **vector** form using the **property**  $\mathbf{z}^T \mathbf{z} = \sum_i z_i^2$

$$J(\mathbf{w}) = \frac{1}{2m} (\mathbf{X}\vec{\mathbf{w}} - \vec{\mathbf{y}})^T (\mathbf{X}\vec{\mathbf{w}} - \vec{\mathbf{y}})$$

# NORMAL EQUATIONS

## DERIVATION OF EQUATIONS



We have to **obtain** the **derivatives** of  $J(w)$  with respect to the vector  $w$ .

$$\nabla_w J(w) = \nabla_w \frac{1}{2m} (X\vec{w} - \vec{y})^T (X\vec{w} - \vec{y})$$

The **gradient** **calculates** all the **derivatives** with **respect** to **all** the weights  $w_j$  **simultaneously**.

Also, because we **represent** the **training dataset** in **matrix form**, we can **calculate** the **gradient** for **all training data**.



# NORMAL EQUATIONS

## DERIVATION OF EQUATIONS



Therefore, we obtain:


$$\nabla_w J(\mathbf{w}) = \nabla_w \frac{1}{2m} (\mathbf{X}\vec{\mathbf{w}} - \vec{\mathbf{y}})^T (\mathbf{X}\vec{\mathbf{w}} - \vec{\mathbf{y}})$$

DERIVE THE EQUALITY

$$\nabla_w J(\mathbf{w}) = \frac{1}{m} (\mathbf{X}^T \mathbf{X}\vec{\mathbf{w}} - \mathbf{X}^T \vec{\mathbf{y}})$$

# NORMAL EQUATIONS

## DERIVATION OF EQUATIONS



**Equating the derivative to zero** to find the **minimum**, we obtain the **vector of weights** that gives this **minimum**:

$$\nabla_w J(w) = \frac{1}{m} (X^T X \vec{w} - X^T \vec{y}) = 0$$

$$\vec{w} = (X^T X)^{-1} X^T \vec{y}$$



# PROBABILISTIC INTERPRETATION

ASSUMPTIONS FOR ERRORS  
LIKELIHOOD FUNCTION

# PROBABILISTIC INTERPRETATION

## M O T I V A T I O N



**WHY DO WE MINIMIZE THE MEAN QUADRATIC  
ERROR FUNCTION AND NOT ANOTHER FUNCTION?**

# PROBABILISTIC INTERPRETATION

## ERROR ASSUMPTIONS



The **output variables**  $y^{(i)}$  can be defined as the **hypothesis**  $h_w$  **plus** an **estimation error**  $\epsilon^{(i)}$  that captures the **effects** that we **did not consider** in our **hypothesis** or contemplates the **random errors**.

$$y^{(i)} = h_w(x^{(i)}) + \epsilon^{(i)}$$

# PROBABILISTIC INTERPRETATION

## ERROR ASSUMPTIONS



The **following assumptions** are made for the **errors**  $\varepsilon^{(i)}$  (random sampling): **IID**

Independent Errors - The probability of one error does not affect the probability of other errors.

Identically distributed errors: errors are sampled from the same probability distribution.

Distributed by a Gaussian probability distribution with mean  $\mu = 0$  and variance  $\sigma^2$ .

$$\varepsilon^{(i)} \sim N(0, \sigma^2)$$

# PROBABILISTIC INTERPRETATION

## ERROR ASSUMPTIONS



Therefore, the **probability density function** of  $\varepsilon^{(i)}$  is given by:

$$p(\varepsilon^{(i)}) = \frac{1}{\sqrt{2\pi}\sigma} e^{\left(-\frac{(\varepsilon^{(i)})^2}{2\sigma^2}\right)}$$

# PROBABILISTIC INTERPRETATION

## ERROR ASSUMPTIONS




**WHY PROPOSE A MODEL THAT IS NORMALLY  
DISTRIBUTED?**



# PROBABILISTIC INTERPRETATION

## ERROR ASSUMPTIONS



Furthermore, the **outputs**  $y^{(i)}$  and the **inputs**  $x^{(i)}$  are set as **random variables**.

Consequently, the following question is posed:


**Given my input data set  $X$  and my weight vector  $\vec{w}$ ,  
what is the probability distribution that models my  
outputs  $\vec{y}$ ?**

# PROBABILISTIC INTERPRETATION

## ERROR ASSUMPTIONS



Following the **assumptions** from the **normal distribution of errors**, it naturally occurs that the **distribution** of the outputs  $y^{(i)}$  follows a **normal form** with mean  $\mathbf{h}_{\mathbf{w}}(\mathbf{x}^{(i)})$  and variance  $\sigma^2$ .



# PROBABILISTIC INTERPRETATION

## LIKELIHOOD FUNCTION



Formally it would look like this:

$$p(y^{(i)} / x^{(i)}; w) = \frac{1}{\sqrt{2\pi}\sigma} e^{\left(-\frac{(y^{(i)} - h_w(x^{(i)}))^2}{2\sigma^2}\right)}$$

$$p(y^{(i)} / x^{(i)}; w) = \frac{1}{\sqrt{2\pi}\sigma} e^{\left(-\frac{(y^{(i)} - w^T x^{(i)})^2}{2\sigma^2}\right)}$$

# PROBABILISTIC INTERPRETATION

## LIKELIHOOD FUNCTION



We define the **conditional probability** for **all training data** (assuming the **samples** were **collected independently**):

$$p(\vec{y}/X; \mathbf{w}) = \prod_{i=1}^m p(y^{(i)}/x^{(i)}; \mathbf{w}) = \prod_{i=1}^m \frac{1}{\sqrt{2\pi}\sigma} e^{\left(-\frac{(y^{(i)} - \mathbf{w}^T x^{(i)})^2}{2\sigma^2}\right)}$$

# PROBABILISTIC INTERPRETATION

## LIKELIHOOD FUNCTION



This equation is called the **likelihood function**, because it describes the **probability** that our **beliefs** about the **reality** of the **observations** (data) are **true**. In other words, that the **hypothesis** we proposed is **correct**.

$$p(\vec{y}/X; \vec{w}) = \prod_{i=1}^m p(y^{(i)}/x^{(i)}; \vec{w}) = \prod_{i=1}^m \frac{1}{\sqrt{2\pi}\sigma} e^{\left(-\frac{(y^{(i)} - \vec{w}^T x^{(i)})^2}{2\sigma^2}\right)}$$

$$p(\mathbf{w}|\mathcal{D}) = \frac{p(\mathcal{D}|\mathbf{w})p(\mathbf{w})}{p(\mathcal{D})}$$

# PROBABILISTIC INTERPRETATION

## LIKELIHOOD FUNCTION




Therefore, we want to **maximize** the **probability** that **our beliefs** about how the **data is distributed** (in a **normal** way).

That is, we want to **maximize** the **likelihood function**. From a **frequentist** perspective (the value of  $\vec{w}$  is **not random**), we want to **find** the **value** of  $\vec{w}$  that **maximizes** the **probability** that the **observations** made of **reality** are **true**.

$$\arg \max_{\vec{w}} L(\vec{w}; X, \vec{y}) = \arg \max_{\vec{w}} \prod_{i=1}^m \frac{1}{\sqrt{2\pi}\sigma} e^{\left(-\frac{(y^{(i)} - w^T x^{(i)})^2}{2\sigma^2}\right)}$$

# PROBABILISTIC INTERPRETATION

## LOG - LIKELIHOOD




**Maximizing a summation** is much easier than doing so with a **multiplication**

$$\arg \max_{\vec{w}} \log(L(\vec{w})) = \arg \max_{\vec{w}} \log\left(\prod_{i=1}^m p(y^{(i)} / x^{(i)}; \vec{w})\right)$$

$$\arg \max_{\vec{w}} \log(L(\vec{w})) = \arg \max_{\vec{w}} \sum_{i=1}^m \log(p(y^{(i)} / x^{(i)}; \vec{w}))$$

# INTERPRETACIÓN PROBABILÍSTICA

## LOG - LOS S



The **maximization problem** is **converted** into a **minimization** one by **scaling** the function with a **minus sign**.

This **function**  $-\log(L(\vec{w}))$  is called **logarithmic loss**.


$$\arg \min_{\vec{w}} -\log(L(\vec{w})) = \arg \min_{\vec{w}} - \sum_{i=1}^m \log(p(y^{(i)} / x^{(i)}; \vec{w}))$$

$$\arg \min_{\vec{w}} -\log(L(\vec{w})) = \arg \min_{\vec{w}} - \sum_{i=1}^m \log\left(\frac{1}{\sqrt{2\pi}\sigma} e^{\left(-\frac{(y^{(i)} - w^T x^{(i)})^2}{2\sigma^2}\right)}\right)$$



# INTERPRETACIÓN PROBABILÍSTICA

## LOG - LOS S



We develop the equation:

$$\arg \min_{\vec{w}} -\log(L(\vec{w})) = \arg \min_{\vec{w}} - \sum_{i=1}^m \log \left( \frac{1}{\sqrt{2\pi}\sigma} \right) + \log \left( e^{-\frac{(y^{(i)} - w^T x^{(i)})^2}{2\sigma^2}} \right)$$

$$\arg \min_{\vec{w}} -\log(L(\vec{w})) = \arg \min_{\vec{w}} - m \log \left( \frac{1}{\sqrt{2\pi}\sigma} \right) - \sum_{i=1}^m -\frac{(y^{(i)} - w^T x^{(i)})^2}{2\sigma^2}$$

# INTERPRETACIÓN PROBABILÍSTICA

## LOG - LOS S



We develop the equation:

$$\arg \min_{\vec{w}} -\log(L(\vec{w})) = \arg \min_{\vec{w}} c + \sum_{i=1}^m \frac{(y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)})^2}{2\sigma^2}$$

$$\arg \min_{\vec{w}} -\log(L(\vec{w})) = \arg \min_{\vec{w}} \frac{1}{2\sigma^2} \sum_{i=1}^m (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)})^2$$

# INTERPRETACIÓN PROBABILÍSTICA

## MEAN SQUARED ERROR



Minimizing the logarithmic loss is the same as minimizing the mean square error  $MSE = J(\vec{w})$ :

$$\arg \min_{\vec{w}} MSE = \arg \min_{\vec{w}} \frac{1}{2m} \sum_{i=1}^m (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)})^2$$

$$\arg \min_{\vec{w}} -\log(L(\vec{w})) = \arg \min_{\vec{w}} \frac{1}{2\sigma^2} \sum_{i=1}^m (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)})^2$$



# AI


**BASIS FUNCTIONS**

# B A S I S   F U N C T I O N S

## L I N E A R I T Y   P R O B L E M



The **real world** has **non-linear** behaviors, so the **linear regression models** seen so far have their **limitations**.




# BASIS FUNCTIONS


## LINEARITY PROBLEM



We can **establish** our **hypothesis** as a **polynomial model**.



$$y = \theta_0 + \theta_1 x + \theta_2 x^2$$



$$y = \sum_{j=0}^5 \theta_j x^j$$

# B A S I S   F U N C T I O N S



## INTRODUCTION TO BASIS FUNCTIONS

If we take this kind of reasoning further, we can **build** models as **linear combinations** of **nonlinear functions**  $\phi_j(x)$ , called **base functions**. Where  $\phi: \mathbb{R}^n \rightarrow \mathbb{R}^k$

$$h_w(x^{(i)}) = \sum_{j=1}^k w_j \phi_j(x^{(i)})$$

$$h_w(x^{(i)}) = w^T \phi(x^{(i)})$$

**NOTA**

$$w \in \mathbb{R}^k$$

$$x^{(i)} \in \mathbb{R}^n$$

Thus we can **build** a **non-linear model**, which is still **parametrized** by **linear weights**  $w$ .

# B A S I S F U N C T I O N S

## AN INTRODUCTION TO BASIS FUNCTIONS



In a more detailed fashion, we have that:

$$\boldsymbol{\phi}(\boldsymbol{x}^{(i)}) = \begin{bmatrix} \phi_1(\boldsymbol{x}^{(i)}) \\ \phi_j(\boldsymbol{x}^{(i)}) \\ \vdots \\ \phi_k(\boldsymbol{x}^{(i)}) \end{bmatrix}$$



# B A S I S F U N C T I O N S

## CLASSICAL LINEAR REGRESSION



In the case of **classical linear regression** we have for a **single** training data  $\phi(x) \in \mathbb{R}^k$  in this case  $k = n$ :

$$\phi_j(x^{(i)}) = x^{(i)} \quad x^{(i)} = \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} \quad w = \begin{bmatrix} w_1 \\ \vdots \\ w_k \end{bmatrix}$$

# B A S I S F U N C T I O N S

## CLASSICAL LINEAR REGRESSION



In the case of **classical linear regression** we have for a **single** training data  $\phi(x) \in \mathbb{R}^k$  in this case  $k = n$ :

$$\phi(x) = \begin{bmatrix} \phi_0(x^{(1)}) \\ \phi_1(x^{(i)}) \end{bmatrix} = \begin{bmatrix} \mathbf{1} \\ x^{(i)} \end{bmatrix} = \begin{bmatrix} \mathbf{1} \\ x_1 \\ \vdots \\ x_n \end{bmatrix}$$

$$h_w(x^{(i)}) = w^T \phi(x^{(i)}) = w_0 + w_1(x_1) + \cdots + w_k(x_n)$$

# B A S I S   F U N C T I O N S

## CLASSICAL LINEAR REGRESSION




In the case of **classical linear regression**, we have for  $m$  training **data**:

$$\phi(x) = \begin{bmatrix} \phi_0(x^{(1)}) & \phi_1(x^{(1)}) \\ \vdots & \vdots \\ \phi_0(x^{(m)}) & \phi_1(x^{(m)}) \end{bmatrix} = \begin{bmatrix} \mathbf{1} & x^{(1)T} \\ \vdots & \vdots \\ \mathbf{1} & x^{(m)T} \end{bmatrix}$$

$$w = [w_0 \quad \dots \quad w_k]$$

# B A S I S   F U N C T I O N S

## CLASSICAL LINEAR REGRESSION




In the case of **classical linear regression**, we have for  $m$  training **data**:

$$\mathbf{h}_w(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}) = \begin{bmatrix} \mathbf{w}_0 + \cdots + \mathbf{w}_k \mathbf{x}_n^{(1)} \\ \vdots \\ \mathbf{w}_0 + \cdots + \mathbf{w}_k \mathbf{x}_n^{(m)} \end{bmatrix}$$

$$\mathbf{h}_w(\mathbf{x}) = \begin{bmatrix} \mathbf{h}_w(\mathbf{x}^{(1)}) \\ \vdots \\ \mathbf{h}_w(\mathbf{x}^{(m)}) \end{bmatrix}$$

# B A S I S   F U N C T I O N S

## D E S I G N   M A T R I X



Therefore, the **design matrix** for any  $\phi(x)$  would be given by:

$$\phi(x) = \begin{bmatrix} \phi_0(x^{(1)}) & \cdots & \phi_{k-1}(x^{(1)}) \\ \vdots & \ddots & \vdots \\ \phi_0(x^{(m)}) & \cdots & \phi_{k-1}(x^{(m)}) \end{bmatrix}$$

$$x^{(i)} \in \mathbb{R}^n$$

Where each row of the matrix  $\phi(x)$  is given by  $\phi_j = \phi(x^{(i)})^T$

# BASIS FUNCTIONS

## LINEAR BASIS FUNCTIONS



Classic linear functions:

$$\phi_j(\mathbf{x}^{(i)}) = \mathbf{x}^{(i)}; \mathbf{x}^{(i)} \in \mathbb{R}^{n+1}$$

$$\mathbf{w} = \begin{bmatrix} \mathbf{w}_0 \\ \vdots \\ \mathbf{w}_k \end{bmatrix}; k = n$$

$$\phi(\mathbf{x}) = \begin{bmatrix} \mathbf{1} & \mathbf{x}^{(1)T} \\ \vdots & \vdots \\ \mathbf{1} & \mathbf{x}^{(m)T} \end{bmatrix} = \begin{bmatrix} \mathbf{1} & \dots & \mathbf{x}_n^{(1)} \\ \vdots & \ddots & \vdots \\ \mathbf{1} & \dots & \mathbf{x}_n^{(m)} \end{bmatrix}$$


$$h_{\mathbf{w}}(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) = \begin{bmatrix} \mathbf{w}_0 + \mathbf{w}_1 \mathbf{x}_1^{(1)} + \dots + \mathbf{w}_k \mathbf{x}_n^{(1)} \\ \vdots \\ \mathbf{w}_0 + \mathbf{w}_1 \mathbf{x}_1^{(m)} + \dots + \mathbf{w}_k \mathbf{x}_n^{(m)} \end{bmatrix}$$

# B A S I S   F U N C T I O N S

## LINEAR BASIS FUNCTIONS



**Classic linear functions:**



# BASIS FUNCTIONS

## POLYNOMIAL BASIS FUNCTIONS



Polynomial basis functions:

$$\phi_j(x) = (x^{(i)})^j; x^{(i)} \in \mathbb{R}^n$$

$$\mathbf{w} = \begin{bmatrix} \mathbf{w}_0 \\ \vdots \\ \mathbf{w}_{k*n+1} \end{bmatrix}$$

$$\phi(x) = \begin{bmatrix} \mathbf{1} & x^{(1)T} & \dots & (x^{(1)T})^{(k-1)} \\ \vdots & \vdots & & \vdots \\ \mathbf{1} & x^{(m)T} & \dots & (x^{(m)T})^{(k-1)} \end{bmatrix}$$

$$\mathbf{h}_w(x) = \mathbf{w}^T \phi(x) = \begin{bmatrix} \mathbf{w}_0 + \mathbf{w}_1 x_1^{(1)} + \dots + \mathbf{w}_{k*n} (x_n^{(1)})^{k-1} \\ \vdots \\ \mathbf{w}_0 + \mathbf{w}_1 x_1^{(m)} + \dots + \mathbf{w}_{k*n} (x_n^{(m)})^{k-1} \end{bmatrix}$$



# BASIS FUNCTIONS

## POLYNOMIAL BASIS FUNCTIONS



**Polynomial basis functions:** concrete example polynomial degree 2 and two training data  $m = 2$

$$\mathbf{x}^{(i)} = \begin{bmatrix} \mathbf{1} \\ x_1 \\ x_2 \end{bmatrix}$$

$$\mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ w_3 \\ w_4 \end{bmatrix}$$

$$\phi(\mathbf{x}) = \begin{bmatrix} \mathbf{1} & x_1^{(1)} & x_2^{(1)} & (x_1^{(1)})^2 & (x_2^{(1)})^2 \\ \mathbf{1} & x_1^{(2)} & x_2^{(2)} & (x_1^{(2)})^2 & (x_2^{(2)})^2 \end{bmatrix}$$

# B A S I S F U N C T I O N S

## POLYNOMIAL BASIS FUNCTIONS



**Polynomial basis functions:** concrete example **polynomial degree 2** and two training data  $m = 2$  with **2 characteristics**.

$$h_w(x) = w^T \phi(x)$$

$$h_w(x) = \begin{bmatrix} \mathbf{1} & x_1^{(1)} & x_2^{(1)} & (x_1^{(1)})^2 & (x_2^{(1)})^2 \\ \mathbf{1} & x_1^{(2)} & x_2^{(2)} & (x_1^{(2)})^2 & (x_2^{(2)})^2 \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ w_3 \\ w_4 \end{bmatrix}$$


$$h_w(x) = \begin{bmatrix} w_0 + w_1 x_1^{(1)} + w_2 x_2^{(1)} + w_3 (x_1^{(1)})^2 + w_4 (x_2^{(1)})^2 \\ w_0 + w_1 x_1^{(2)} + w_2 x_2^{(2)} + w_3 (x_1^{(2)})^2 + w_4 (x_2^{(2)})^2 \end{bmatrix}$$

# B A S I S   F U N C T I O N S

## POLYNOMIAL BASIS FUNCTIONS



**Polynomial basis functions:**



# B A S I S F U N C T I O N S


## F O U R I E R S E R I E S



**Fourier series**

$$\phi_0(x) = 1$$

$$\phi_j(x) = \cos(\varpi_j x^{(i)} + \psi_j); j > 0$$




# B A S I S F U N C T I O N S

## R A D I A L B A S I S



Radial basis function:

$$\phi_j(x) = e^{-\frac{1}{2l} \sum_{r=1}^n (x_r^{(i)} - \mu_{j,r})^2}$$



# B A S I S   F U N C T I O N S


## S I G M O I D   F U N C T I O N



Sigmoid function:

$$\phi_j(x^{(i)}) = \sigma\left(\frac{x^{(i)} - \mu_j}{s}\right)$$

$$\sigma(a) = \frac{1}{1 + e^{-a}}$$



# B A S I S   F U N C T I O N S

## M A X I M U M   L I K E L I H O O D



The **same** estimation **model** is proposed **plus** an **error**:


$$y^{(i)} = h_w(x^{(i)}) + \varepsilon^{(i)}$$

Where:

$$h_w(x^{(i)}) = w^T \phi(x)$$

# B A S I S   F U N C T I O N S

## M A X I M U M   L I K E L I H O O D




**Assuming** again that the **error** is **normally distributed** with **mean 0** and **variance  $\beta^{-1}$** . It is defined for  **$m$  training data**:

$$p(\vec{y}/X, \mathbf{w}, \beta) = \prod_{i=1}^m p(y^{(i)}/x^{(i)}; \mathbf{w}) = \prod_{i=1}^m \sqrt{\frac{\beta}{2\pi}} e^{\left(-\frac{\beta}{2}(y^{(i)} - \mathbf{w}^T \phi(x^{(i)}))^2\right)}$$



# B A S I S F U N C T I O N S

## M A X I M U M L I K E L I H O O D



Since the **likelihood function** is **parameterized** in  $w$  and  $\beta$  the logarithmic likelihood is written like this:

$$\log p(\vec{y}/w, \beta) = \log \prod_{i=1}^m \sqrt{\frac{\beta}{2\pi}} e^{\left(-\frac{\beta}{2}(y^{(i)} - w^T \phi(x^{(i)}))^2\right)}$$

$$\log p(\vec{y}/w, \beta) = \sum_{i=1}^m \log \sqrt{\frac{\beta}{2\pi}} e^{\left(-\frac{\beta}{2}(y^{(i)} - w^T \phi(x^{(i)}))^2\right)}$$

# B A S I S   F U N C T I O N S

## M A X I M U M   L I K E L I H O O D




Developing:

$$\log p(\vec{y}/w, \beta) = \sum_{i=1}^m \log \sqrt{\frac{\beta}{2\pi}} + \log e^{\left(-\frac{\beta}{2}(y^{(i)} - w^T \phi(x^{(i)}))^2\right)}$$

$$\log p(\vec{y}/w, \beta) = m \log \sqrt{\frac{\beta}{2\pi}} - \sum_{i=1}^m \frac{\beta}{2} (y^{(i)} - w^T \phi(x^{(i)}))^2$$

# B A S I S F U N C T I O N S

## M A X I M U M L I K E L I H O O D



Developing:


$$\log p(\vec{y}/w, \beta) = \frac{m}{2} \log \beta - \frac{m}{2} \log 2\pi - \sum_{i=1}^m \frac{\beta}{2} \left( y^{(i)} - w^T \phi(x^{(i)}) \right)^2$$

$$\log p(\vec{y}/w, \beta) = \frac{m}{2} \log \beta - \frac{m}{2} \log 2\pi - \beta E_D(w)$$

$$E_D(w) = \frac{1}{2} \left( y^{(i)} - w^T \phi(x^{(i)}) \right)^2$$

# B A S I S F U N C T I O N S

## M A X I M U M L I K E L I H O O D




Developing:

$$-\log p(\vec{y}/w, \beta) = -\frac{m}{2} \log \beta + \frac{m}{2} \log 2\pi + \beta E_D(w)$$

$$\arg \min_{\vec{w}} -\log(L(\vec{w})) = \arg \min_{\vec{w}} \frac{\beta}{2} \sum_{i=1}^m \left( y^{(i)} - w^T \phi(x^{(i)}) \right)^2$$

# B A S I S   F U N C T I O N S

## N O R M A L   E Q U A T I O N S



Calculating the **gradient**, **setting** it equal to **zero** and solving for the **vector  $w$** :

$$\mathbf{w} = (\boldsymbol{\phi}^T \boldsymbol{\phi})^{-1} \boldsymbol{\phi}^T \vec{\mathbf{y}}$$