



MACHINE LEARNING

REINFORCEMENT LEARNING

AGENDA

- 01** Introduction
- 02** Markov Decision Processes
- 03** Policy and Value Function
- 04** Value Iteration and Policy Iteration
- 05** Learning MDP Models
- 06** Continuous state MDPs



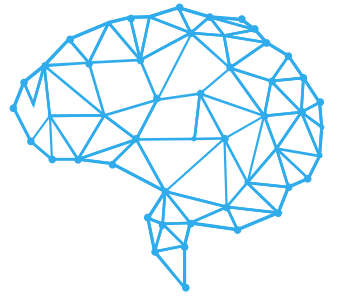


AI

INTRODUCTION

REINFORCEMENT LEARNING

INTRODUCTION



There are many **problems**, where it is **very difficult** to **provide** the “**right**” **answers** to the algorithm, as we have been doing for supervised learning.

For example, if we have just built a **four-legged robot** and are **trying** to **program** it to **walk**, then initially we **have no idea** what the “**correct**” **actions** to take are to make it walk (**credit assignment problem**).

In **reinforcement learning**, we will **provide algorithms** only a **reward function**, to **indicate** when the **learning agent** is **doing well** or **poorly**.



It will be the learning algorithm's job to figure out how to choose actions over time to maximize rewards

REINFORCEMENT LEARNING

INTRODUCTION



Many of the applications of reinforcement learning include:

- **Autonomous helicopter flight.**
- **Robot Legged Locomotion.**
- **Cell-phone network routing.**
- **Marketing strategy selection.**
- **Factory control.**
- **Efficient web indexing.**
- **Traffic light control.**
- **Chemical reaction optimization.**
- **Personalized recommendations.**
- **Bidding and Advertisement.**





AI

MARKOV DECISION PROCESSES

REINFORCEMENT LEARNING

MARKOV DECISION PROCESSES



A **Markov decision process** is a tuple:

$$(S, A, \{P_{sa}\}, \gamma, R)$$

- S : the set of **states** (i.e. all positions and orientations of a helicopter).
- A : the set of **actions** (i.e. all directions in which you can push the helicopter's sticks).
- P_{sa} : **state transition probability distributions**. For **each state** $s \in S$ and action $a \in A$, P_{sa} is a **distribution** over the **state space**. Thus, P_{sa} **gives** the **distribution** over what **states** the **algorithm** will **transition**, if it takes **action** a in **state** s .

$$\sum_{s'} P_{sa}(s') = 1$$

REINFORCEMENT LEARNING

MARKOV DECISION PROCESSES



A **Markov decision process** is a tuple:

$$(S, A, \{P_{sa}\}, \gamma, R)$$

- $\gamma \in [0, 1)$: is called the **discount factor**.
- $R: S \times A \rightarrow \mathbb{R}$: the **reward function**. (It can be written only as a function of S , $R: S \rightarrow \mathbb{R}$).

REINFORCEMENT LEARNING

MARKOV DECISION PROCESSES



EXAMPLE:

We will **define** a **robot navigation task**, in which you **have** a **robot** in a **grid environment** where **all gray cells** are **obstacles**.

The **reward** is **defined** by the **+1** and the **punishment** as **-1**.


The **robot** can move in **all cells** **except** for the **gray cells**.

11 states

			+1
			-1

STOCHASTIC ACTIONS

$$A = \{N, W, S, E\}$$



REINFORCEMENT LEARNING

MARKOV DECISION PROCESSES



EXAMPLE:

The **robot** can transit from one **state** to another **state** in **single steps**.

When the robot reaches (4,3) or (4,2) the world ends.

(1, 3)	(2, 3)	(3, 3)	(4, 3) +1
(1, 2)	(2, 2)	(3, 2)	(4, 2) -1
(1, 1)	(2, 1)	(3, 1)	(4, 1)

Diagram showing transitions from state (3, 1):

- 80% probability to (3, 2)
- 10% probability to (3, 3)
- 10% probability to (3, 1)

Transitions

$$P_{s_t a_t}(s')$$

$$P_{(3,1)N}((3, 2)) = 0.8$$

$$P_{(3,1)N}((4, 1)) = 0.1$$

$$P_{(3,1)N}((2, 1)) = 0.1$$

$$P_{(3,1)N}((3, 3)) = 0$$

⋮

Rewards

$$R(s)$$

$$R((4, 3)) = +1$$

$$R((4, 2)) = -1$$


$$R(s) = -0.2$$

(fuel consumption)

⋮

REINFORCEMENT LEARNING

MARKOV DECISION PROCESSES



The dynamics of a **Markov decision process** is as follows:

1. The **algorithm starts** in a **state** s_0 and **choose** an **action** $a_0 \in A$.
2. There is a **random transition** from **state** s_0 to **state** s_1 **drawn according** to $s_1 \sim P_{s_0 a_0}$.
3. The **process repeats indefinitely** where there is a **random selection** to **traverse** from state s_t to state s_{t+1} , which is **drawn from** $s_{t+1} \sim P_{s_t a_t}$.

$$s_0 \xrightarrow{a_0} s_1 \xrightarrow{a_1} s_2 \xrightarrow{a_2} \dots$$


4. The **total payoff** is be given by:

$$R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

(Rewards diminish as times goes by).

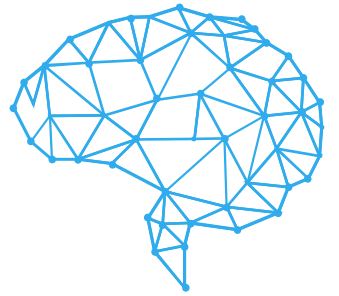
REINFORCEMENT LEARNING

MARKOV DECISION PROCESSES



REINFORCEMENT LEARNING

MARKOV DECISION PROCESSES



The **objective** in reinforcement learning is to **choose actions over time** to **maximize** the **expected value** of the **total payoff**:

$$E[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots]$$

The **reward** at **timestep** t is **discounted** by a **factor** of γ^t where $\gamma \in [0, 1)$ is the **discount factor**.

Therefore, to **make** this **expectation large**, we would like to:

- **Obtain positive rewards as soon as possible.**
- **Postpone negative rewards, as long as possible.**



AI

POLICY AND VALUE FUNCTION

REINFORCEMENT LEARNING

POLICY AND VALUE FUNCTION



POLICY

A **policy** is any function $\pi : S \rightarrow A$ that maps from states to actions (for every state, what action is recommended to take in that state).


EXAMPLE: (Optimal Policy)

(1, 3) →	(2, 3) →	(3, 3) →	+1
(1, 2) ↑	(2, 2)	(3, 2) ↑	-1
(1, 1) ↑	(2, 1) ←	(3, 1) ←	(4, 1) ←

MDPs are very good in finding the best tradeoffs between actions (policies).

REINFORCEMENT LEARNING

POLICY AND VALUE FUNCTION



VALUE FUNCTION:

For **any policy** π , we **define** the **value function** $V^\pi: S \rightarrow \mathbb{R}$ as the **expected total payoff** upon **starting** in **state** s and **taking actions according** to the **policy** π .

$$V^\pi(s) = E[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots | s_0 = s, \pi]$$

As a **side note**, this **notation** is **technically incorrect** because we **cannot condition** on π , which **does not represent** a random variable.

REINFORCEMENT LEARNING

POLICY AND VALUE FUNCTION



EXAMPLE: Value Function and Policy

Very bad policy π

(1, 3) →	(2, 3) →	(3, 3) →	+1
(1, 2) ↓	(2, 2)	(3, 2) →	-1
(1, 1) →	(2, 1) →	(3, 1) ↑	(4, 1) ↑

Value Function $V^\pi(s)$


$V^\pi((1, 3))$ 0.52	$V^\pi((2, 3))$ 0.73	$V^\pi((3, 3))$ 0.77	+1
$V^\pi((1, 2))$ -0.90		$V^\pi((3, 2))$ -0.82	-1
$V^\pi((1, 1))$ -0.88	$V^\pi((2, 1))$ -0.87	$V^\pi((3, 1))$ -0.85	$V^\pi((4, 1))$ -1

$$V^\pi(s) = E[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots | s_0 = s, \pi]$$

Better to start at the top

REINFORCEMENT LEARNING

POLICY AND VALUE FUNCTION



BELLMAN EQUATIONS:

We can **represent** the **value function** in the **following way**:

$$V^{\pi}(s) = E[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots | s_0 = s, \pi]$$

$$V^{\pi}(s) = E[R(s_0) + \gamma(R(s_1) + \gamma R(s_2) + \dots) | s_0 = s, \pi]$$

Immediate reward
for starting in state s

Expected sum of future
discounted rewards.

$V^{\pi}(s_1)$

Therefore, we can **write** the **equation** as a **recursive form of itself** ($s_0 \rightarrow s$ and $s_0 \rightarrow s'$) **using** the **definition of expected value** $E[X] = \sum x_i p_i$ (s' is a **random variable**):

$$V^{\pi}(s) = R(s) + \gamma \sum_{s' \in \mathcal{S}} P_{s\pi(s)}(s') V^{\pi}(s')$$

REINFORCEMENT LEARNING

POLICY AND VALUE FUNCTION



BELLMAN EQUATIONS:

Given a fixed policy π , its value function V^π satisfies Bellman equations:

$$V^\pi(s) = R(s) + \gamma \sum_{s' \in \mathcal{S}} P_{s\pi(s)}(s') V^\pi(s')$$

The **first term** is the **immediate reward** for **starting** in **state s** .

The **second term** gives the **expected sum of discounted rewards** obtained after the **first step** in the **MDP**.

REINFORCEMENT LEARNING

POLICY AND VALUE FUNCTION



EXAMPLE: Bellman Equation

Let us **assume** that we **start** at **state** (3,1) and the **policy** π **indicates** to take a **step** in the **north direction** $\pi((3, 1)) = N$.

The **Bellman Equation** is:

$$V^\pi(s) = R(s) + \gamma \sum_{s' \in S} P_{s\pi(s)}(s') V^\pi(s')$$

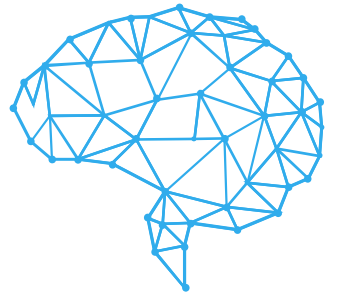
$$V^\pi((3, 1)) = R((3, 1)) + \gamma[0.8V^\pi((3, 2)) + 0.1V^\pi((4, 1)) + 0.1V^\pi((2, 1))]$$

You will have **11 variables** and **11 equations**, each one **corresponding** to each **state**.

(1, 3) →	(2, 3) →	(3, 3) →	+1
(1, 2) ↓	(2, 2)	(3, 2) →	-1
(1, 1) →	(2, 1) →	(3, 1) ↑	(4, 1) ↑

REINFORCEMENT LEARNING

POLICY AND VALUE FUNCTION



OPTIMAL VALUE FUNCTION:

The **optimal value function** will be the **best possible expected sum of discounted rewards** that can be attained **using any policy**.

$$V^*(s) = \max_{\pi} V^{\pi}(s)$$

The **Bellman representation** will be:

$$V^*(s) = R(s) + \max_{a \in A} \gamma \sum_{s' \in S} P_{sa}(s') V^*(s')$$

Expected future rewards if we take the **BEST** action a in state s .

The **first term** above is the **immediate reward**. The **second term** is the **maximum over all actions a** of the **expected future sum of discounted rewards** we'll get upon after action a .

REINFORCEMENT LEARNING

POLICY AND VALUE FUNCTION



OPTIMAL POLICY:

The **optimal policy** π^* indicates the **best action** a to **take** in **state** s , which **maximizes** the **expected future rewards**.

$$\pi^*(s) = \arg \max_{a \in A} \sum_{s' \in S} P_{sa}(s') V^*(s')$$

This **derives** in the **following expression**:

$$V^*(s) = V^{\pi^*}(s) \geq V^{\pi}(s)$$

In conclusion, π^* is the **optimal policy** for all states s .

We can use the same policy π^* no matter what the initial state of our MDP is.

REINFORCEMENT LEARNING

POLICY AND VALUE FUNCTION



OPTIMAL POLICY:

To find the **optimal policy** $\pi^*(s)$, we would **need** to **obtain** the **best value function** $V^*(s')$ and **substitute it** in the **equation** we have just obtained:

$$\pi^*(s) = \arg \max_{a \in A} \sum_{s' \in S} P_{sa}(s') V^*(s')$$

The **problem** is **finding** $V^*(s')$ which corresponds to:

$$V^*(s) = \max_{\pi} V^{\pi}(s)$$

The **bottleneck** is that you would **need** to **solve** an **exponentially large number** of **systems** of **equations** **because** there is an **exponential number** of **policies** π .

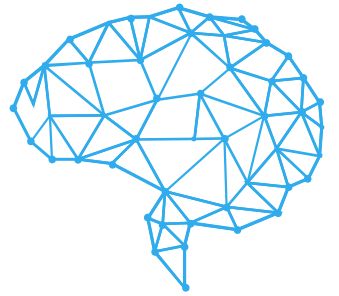


AI

**VALUE ITERATION AND
POLICY ITERATION**

REINFORCEMENT LEARNING

VALUE ITERATION AND POLICY ITERATION



VALUE ITERATION:

1. For each state s , initialize $V(s) := 0 \quad \forall s$.

2. Repeat until convergence{

For every state, update $V(s) := R(s) + \max_{a \in A} \gamma \sum_{s' \in S} P_{sa}(s') V(s')$.

}


3. Substitute $V^*(s)$ in $\pi^*(s) = \arg \max_{a \in A} \sum_{s' \in S} P_{sa}(s') V^*(s')$.

The algorithm repeatedly tries to update the estimated value function using Bellman Equations. By implementing this you ensure that the value function converges to the optimum.

$$V(s) \rightarrow V^*(s)$$

REINFORCEMENT LEARNING

VALUE ITERATION AND POLICY ITERATION



VALUE ITERATION:

There are two possible ways of performing the updates in the inner loop.

$$V(s) := R(s) + \max_{a \in A} \gamma \sum_{s' \in S} P_{sa}(s') V(s').$$

1. **Synchronous update** (current estimate \rightarrow new estimate):
 - **Compute** the **new values** of $V(s)$.
 - **Overwrite** the **old values** with the **new values**.
2. **Asynchronous updates**
 - **Loop over** the **states** (**predefined order**).
 - **Update** the **values** of $V(s)$ **one** at a **time**.

REINFORCEMENT LEARNING

VALUE ITERATION AND POLICY ITERATION



EXAMPLE: Value Iteration

Running value iteration on our previous example, we obtain the following results for $V^*(s)$:

$V^*(s)$			
$V^\pi((1, 3))$ 0.86	$V^\pi((2, 3))$ 0.90	$V^\pi((3, 3))$ 0.93	+1
$V^\pi((1, 2))$ 0.82		$V^\pi((3, 2))$ 0.69	-1
$V^\pi((1, 1))$ 0.78	$V^\pi((2, 1))$ 0.75	$V^\pi((3, 1))$ 0.71	$V^\pi((4, 1))$ 0.49

REINFORCEMENT LEARNING

VALUE ITERATION AND POLICY ITERATION



EXAMPLE: Value Iteration

Let us see how we **obtain** $\pi^*(s)$ for a **single state**: (3, 1)

$$W = \sum_{s' \in S} P_{sa}(s') V(s') = (0.8 * 0.75) + (0.1 * 0.69) + (0.1 * 0.71) = 0.78$$

$$N = \sum_{s' \in S} P_{sa}(s') V(s') = (0.8 * 0.69) + (0.1 * 0.75) + (0.1 * 0.49) = 0.67$$

$$S = \sum_{s' \in S} P_{sa}(s') V(s') = (0.8 * 0.71) + (0.1 * 0.75) + (0.1 * 0.49) = 0.69$$

$$E = \sum_{s' \in S} P_{sa}(s') V(s') = (0.8 * 0.49) + (0.1 * 0.69) + (0.1 * 0.71) = 0.53$$

The action west maximizes our future rewards if we start at state (3,1).

$V^*(s)$

$V^\pi((1,3))$ 0.86	$V^\pi((2,3))$ 0.90	$V^\pi((3,3))$ 0.93	+1
$V^\pi((1,2))$ 0.82		$V^\pi((3,2))$ 0.69	-1
$V^\pi((1,1))$ 0.78	$V^\pi((2,1))$ 0.75	$V^\pi((3,1))$ 0.71	$V^\pi((4,1))$ 0.49

REINFORCEMENT LEARNING

VALUE ITERATION AND POLICY ITERATION



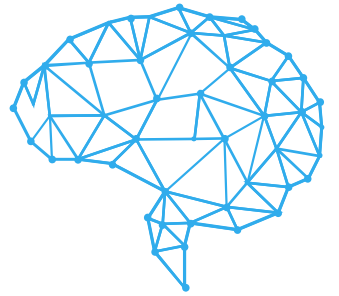
EXAMPLE: Value Iteration

Applying the **previous procedure** to **all states**, we **obtain** our **optimal policy** (best actions $a = \pi(s)$ to take **at every state s**).

$\pi^*(s)$			
(1, 3) →	(2, 3) →	(3, 3) →	+1
(1, 2) ↑	(2, 2)	(3, 2) ↑	-1
(1, 1) ↑	(2, 1) ←	(3, 1) ←	(4, 1) ←

REINFORCEMENT LEARNING

VALUE ITERATION AND POLICY ITERATION



POLICY ITERATION:

1. Initialize π randomly.
2. Repeat until convergence{
 - a) Let $V := V^\pi$ (Solve Bellman equations using policy π).
 - b) For every state, compute $\pi(s) := \underset{a \in A}{\operatorname{argmax}} \sum_{s' \in S} P_{sa}(s') V(s')$.}

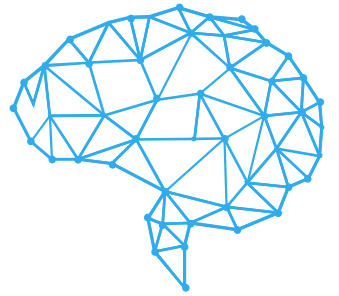
The **policy** π found in **step (b)** is also **called** the **policy** that is **greedy** with **respect** to V .
By **implementing this** you **ensure** that the **value function converges** to the **optimum**.

$$V(s) \rightarrow V^*(s)$$

$$\pi(s) \rightarrow \pi^*(s)$$

REINFORCEMENT LEARNING

VALUE ITERATION AND POLICY ITERATION



POLICY ITERATION VS VALUE ITERATION:

POLICY ITERATION	VALUE ITERATION
For small MDPs it is very fast.	For small MDPs it is slower.
For MDPs with large state spaces it is computationally expensive (large matrices)	For MDPs with large state spaces it is has less computational expense.

Value iteration seems to be **used more often** than policy iteration.



AI

**LEARNING MDP
MODELS**

REINFORCEMENT LEARNING

LEARNING MDP MODELS



In **many realistic problems**, we are **not given** state transition probabilities and rewards **explicitly**. Therefore, we need to **estimate them from data**.

More generally, in a Markov Decision Process $(S, A, \{P_{sa}\}, \gamma, R)$ the **states S** , the **actions A** and the **discount factors γ** are **almost always known**. The **discount factor** is **chosen based on how much tradeoff** do you want **between current and future rewards**.

We can **easily derive** the **maximum likelihood estimates** for the **state transition probabilities**:

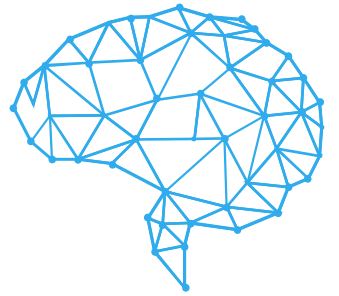
$$P_{sa}(s') = \frac{\text{\# times we took action } a \text{ in state } s \text{ and got to } s'}{\text{\# times we took action } a \text{ in state } s}$$

If there is the **case** that you have **never taken action a in state s** (0/0) we can **compute the uniform distribution over all states** :

$$P_{sa}(s') = \frac{1}{|S|}$$

REINFORCEMENT LEARNING

LEARNING MDP MODELS



Putting together model learning and value iteration, here is **one possible algorithm** for learning in an **MDP** with **unknown state transition probabilities**:

1. Initialize π randomly.
2. Repeat {
 - a) Execute π in the **MDP** for some **number of trials**
 - b) Using the **accumulated experience** in the **MDP**, update our **estimates** for P_{sa} (and R , if applicable).
 - c) Apply value iteration with the **estimated state transition probabilities** and **rewards** to get a **new estimated value function V** (use this update to initialize V in next iteration).
 - (d) Update π to be the **greedy policy** with respect to V .}